

OpenAI Codex CLI Training

AI-Powered Terminal Coding Agent

Press Space for next page →

Contact Info

Ken Kousen Kousen IT, Inc.

- ken.kousen@kousenit.com
- <http://www.kousenit.com>
- <http://kousenit.org> (blog)
- Social Media:
 - [@kenkousen](https://twitter.com/kenkousen) (twitter)
 - [@kenkousen@foojay.social](https://mastodon.foojay.social/@kenkousen) (mastodon)
 - [@kousenit.com](https://bluesky.kousenit.com/@kousenit) (bluesky)
- *Tales from the jar side* (free newsletter)
 - <https://kenkousen.substack.com>
 - <https://youtube.com/@talesfromthejarside>

Course Overview

5-Hour Hands-On Workshop

- Installation and authentication strategies
- Terminal UI and navigation
- Sandbox modes and approval policies
- Real-world coding projects

Topics We'll Cover

- Advanced TOML configuration
- **Agent Skills** - Reusable workflows (NEW!)
- MCP services integration
- Memory with AGENTS.md
- Custom prompts and profiles
- Multi-model provider support

Prerequisites

- Command-line experience
- Basic programming knowledge
- Git familiarity
- Docker (for advanced exercises)

What is OpenAI Codex CLI?

Lightweight Terminal-Based Coding Agent

Key Features

- Multi-model support (GPT, Claude, Ollama)
- Built-in safety with sandbox modes
- Rich configuration system
- Model Context Protocol (MCP) integration

Advanced Capabilities

- Session persistence and resumption
- Custom prompts and profiles
- CI/CD compatible
- Headless execution

Authentication Options

ChatGPT Account (Recommended)

- Uses existing ChatGPT subscription
- Zero Data Retention (ZDR)
- Simplified login flow

API Key

- Direct API access
- Pay-per-use pricing
- More configuration required

Model Support

- **GPT-5.2-Codex** - Latest, default for API (January 2026)
- **GPT-5-Codex** - Stable workhorse model
- **GPT-5-Codex-Mini** - Cost-effective, 4x more usage
- **GPT-5.1-Codex-Max** - Long-running project-scale work
- Anthropic Claude via API
- Local models via Ollama

Installation Methods

```
# NPM (recommended)
npm install -g @openai/codex

# Homebrew (macOS/Linux)
brew install --cask codex

# Direct binary download
# Visit: https://github.com/openai/codex/releases
```

Verify Installation

```
codex --version
```

Configuration Locations

- **Config:** `~/.codex/config.toml`
- **Prompts:** `~/.codex/prompts/`
- **Logs:** `~/.codex/log/`

ChatGPT Account Login

```
# Interactive login  
codex login  
  
# Headless login for servers  
codex login --headless
```

API Key Authentication

```
# Set environment variable
export OPENAI_API_KEY="your-key"

# Or configure in TOML
echo 'api_key = "your-key"' >> ~/.codex/config.toml
```

Verify Authentication

```
codex "Hello, are you working?"
```

Starting Codex

```
# Interactive mode (default)
codex

# With initial prompt
codex "explain this codebase"

# Execute and exit mode
codex exec "generate a README"
```

Key Bindings

- `Enter` - Submit prompt
- `Ctrl+C` - Cancel current operation
- `Ctrl+D` - Exit Codex
- `Tab` - Autocomplete
- `/` - Access slash commands

Slash Commands

- `/status` - Show session info & token usage
- `/diff` - Review all pending changes
- `/clear` - Clear conversation history
- `/save` - Save current session
- `/help` - Show available commands
- `/settings` - Adjust runtime settings

/diff Command - Review Changes

```
--- a/src/main.py
+++ b/src/main.py
@@ -10,7 +10,9 @@ def process_data(input_file):
-    data = json.load(f)
+    with open(input_file, 'r') as f:
+        data = json.load(f)
    return data

3 files changed, 47 insertions(+), 12 deletions(-)
```

Review line-by-line before approving!

/status Command (v0.35+)

Shows comprehensive session information:

```
Current model: gpt-5.2-codex
```

```
Session ID: abc123
```

```
Token usage: 15,432 / 128,000
```

```
Cost estimate: $0.46
```

```
Time elapsed: 12m 34s
```

/review Command (v0.76+)

Launch code review without modifying your working tree:

```
# Review uncommitted changes  
/review  
  
# Review with specific focus  
/review Check for security vulnerabilities  
  
# Review changes against a branch  
/review Compare with main branch
```

Built-in Reviewers

- **Security** - OWASP patterns, injection risks
- **Performance** - N+1 queries, memory leaks
- **Style** - Naming conventions, code structure
- **Tests** - Coverage gaps, edge cases

Reviewers analyze diffs without executing code

Search Your Codebase

```
# Fast text search with ripgrep
rg "TODO"
rg "authenticate"
rg "database connection"
```

- Respects .gitignore and runs fast on large repos
- Pipe matches into Codex for follow-up analysis
- Keep the agent focused by sharing only relevant snippets
- Great starting point for exploration and debugging

Web Search Capabilities

```
# ~/.codex/config.toml  
web_search_request = true
```

- Search the entire web for solutions
- Find latest documentation and tutorials
- Access Stack Overflow answers
- Get real-time information
- Research libraries and frameworks

Using Web Search

```
# In interactive mode, Codex can search the web
codex
> "Search for the latest React 18 features"
> "Find Python async/await best practices"
> "What are the breaking changes in Spring Boot 3?"
```

- Automatic web search when needed
- Current information beyond training cutoff
- Verify solutions with official docs

Image Inputs

Attach screenshots and design specs for visual context:

```
# From command line
codex -i screenshot.png "Explain this error"
codex -i mockup.png "Implement this design"
codex -i diagram.png "Generate code for this architecture"

# Multiple images
codex -i error.png -i logs.png "Debug this issue"
```

Image Input Use Cases

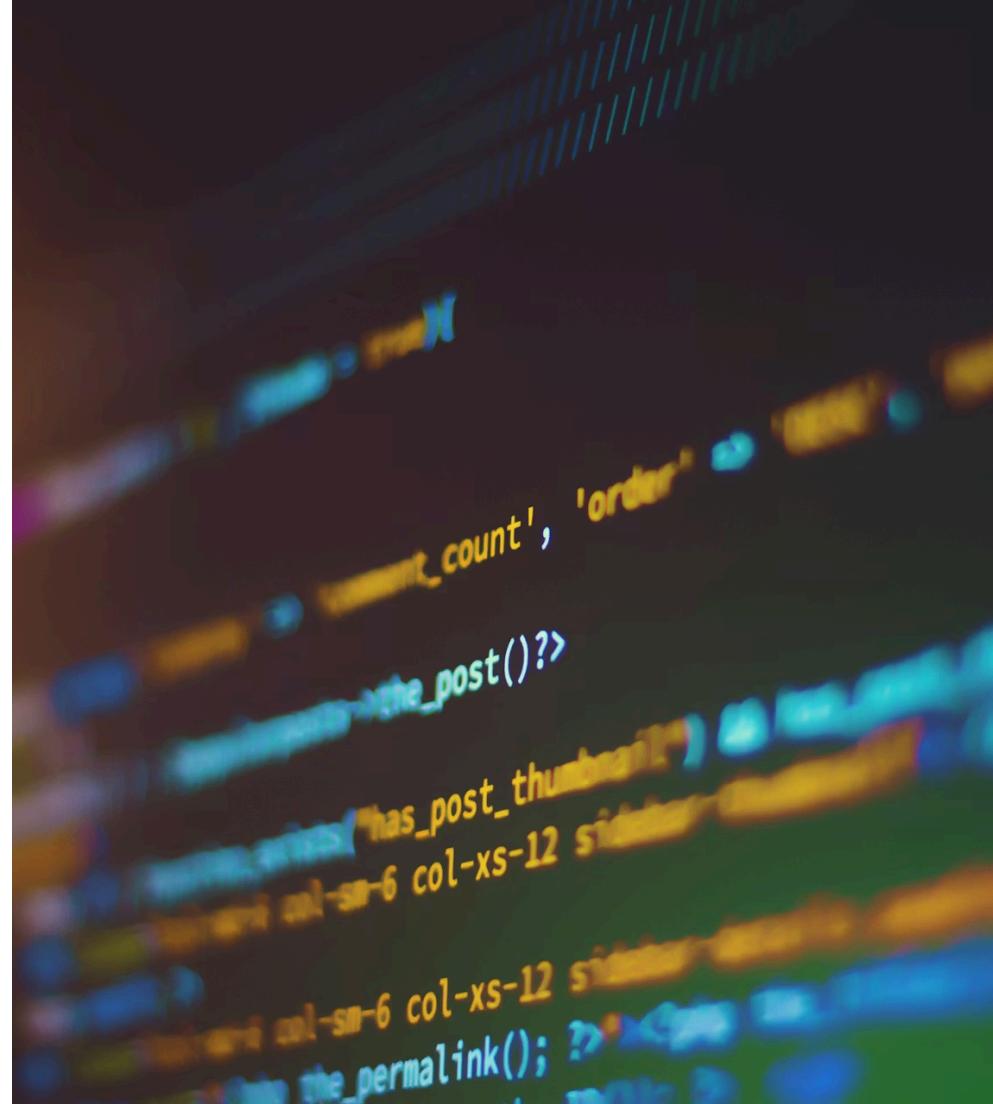
- **Debug UI errors** - Share error dialogs, stack traces
- **Implement designs** - Convert mockups to code
- **Analyze diagrams** - Generate from architecture docs
- **Review screenshots** - Identify accessibility issues
- **Compare outputs** - "Why does this look different?"

Paste images directly in the TUI composer!

Core Features

Essential Capabilities

Master the fundamentals



Sandbox Modes



- **read-only** - No file modifications
- **workspace-write** - Default; writes limited to the workspace
- **danger-full-access** - No sandboxing (use carefully!)

Approval Policies

- **untrusted** - Run only trusted commands without prompting
- **on-request** - Approve risky actions
- **on-failure** - Approve only on failures
- **never** - No approval prompts

Setting Safety Options

```
# Set sandbox mode
codex --sandbox read-only

# Set approval policy
codex --ask-for-approval on-request

# Bypass all safety (dangerous!)
codex --dangerously-bypass-approvals-and-sandbox
```

Project Memory: AGENTS.md

Automatic Context Loading

- Place `AGENTS.md` in project root
- Loaded automatically with first prompt
- Configurable size limit (default: 32KB)

Example AGENTS.md

```
# Project: E-Commerce Platform

## Tech Stack
- Backend: Node.js + Express
- Database: PostgreSQL
- Frontend: React + TypeScript
```

AGENTS.md Best Practices

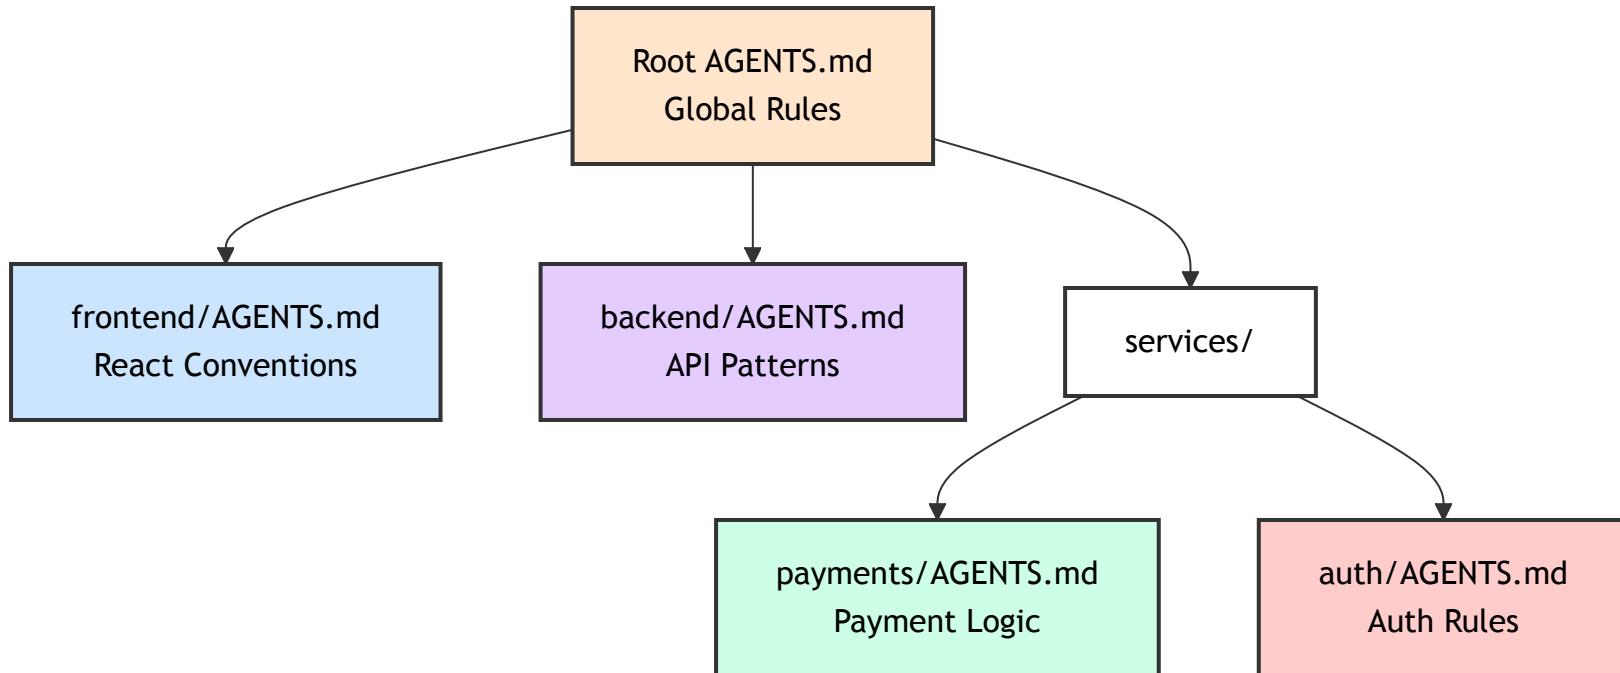
Conventions

- Use async/await for all async operations
- Follow RESTful API patterns
- Write tests for all new features

Current Focus

Working on payment integration with Stripe

Hierarchical AGENTS.md (v0.39+)



Rules cascade: Subfolder overrides parent

Context Cascade Benefits

- Global rules apply everywhere
- Subfolder rules override parent
- Each team owns their conventions
- Frontend/backend stay independent
- Microservices maintain autonomy

Custom Prompts

Creating Custom Prompts

1. Create `.md` file in `~/codex/prompts/`
2. Access via slash commands
3. Reusable across projects

Example Custom Prompt

```
# ~/.codex/prompts/refactor.md
Refactor the selected code following these principles:
1. Extract complex logic into small functions
2. Use meaningful variable names
3. Add appropriate error handling
```

Prompt Library Highlights

- Prebuilt prompts live in `~/ .codex/prompts/` (see repo `prompts/README.md`)
- Core templates: `refactor` , `security-audit` , `test-gen` , `pr-review` , `api-upgrade` , `perf-fix`
- Customize or fork them for your team's workflow and slash commands

```
/refactor  
/security-audit  
/test-gen  
/pr-review  
/api-upgrade  
/perf-fix
```

Prompt Arguments Workaround

- Codex doesn't support `$ARGUMENTS` like Claude Code
- Solution: Use shell scripts as wrappers
- Scripts can accept parameters and build dynamic prompts
- Store in `~/.codex/scripts/` for reuse

Prompt Arguments: Implementation

```
#!/bin/bash
# ~/.codex/scripts/review-file.sh

FILE=$1
FOCUS=$2

cat > /tmp/review-prompt.md << EOF
Review the file ${FILE} focusing on ${FOCUS}:
- Check for bugs and errors
- Suggest improvements
- Rate code quality
EOF


codex exec "$(cat /tmp/review-prompt.md)"
```

Usage: ./review-file.sh UserService.java security

Agent Skills

Reusable Workflows

December 2025



What Are Agent Skills?

- **Reusable instruction bundles** with optional scripts and resources
- **Follows [agentskills.io spec](#)** (same as Claude Code!)
- **Progressive loading:** Only name/description loaded at startup
- **Two invocation modes:** Explicit (`$skill-name`) or implicit (auto-detect)
- **Skills replace complex prompts** for multi-step workflows

Skill Locations

Scope	Location	Use Case
User	<code>~/.codex/skills/</code>	Personal workflows
Repository	<code>.codex/skills/</code>	Team-shared skills
Admin	System-managed	Enterprise policies

Skills load in precedence order: repo → user → admin

Skill Structure

```
my-skill/
├── SKILL.md          # Required: YAML frontmatter + instructions
├── SKILL.toml         # Optional: icons, brand color, defaults
├── scripts/           # Optional: executable code
├── references/        # Optional: documentation
└── assets/            # Optional: templates, resources
```

SKILL.md Format

```
---
```

```
name: security-review
description: >
  Perform security analysis. Use when asked about vulnerabilities,
  security audit, or "is this code secure".
---

# Security Review

Analyze codebases for security vulnerabilities...

## Workflow
1. Reconnaissance - identify tech stack
2. Dependency analysis - check for CVEs
3. Code analysis - scan for patterns
4. Generate report - create SECURITY REVIEW.md
```

Invoking Skills

Explicit Invocation

```
# Use $ prefix to invoke directly  
$skill-creator Create a skill for commit messages  
$create-plan Design a new authentication system
```

Implicit Invocation

```
# Codex auto-selects based on task match  
"Review this code for security vulnerabilities"  
# → Automatically invokes security-review skill if installed
```

Built-in Skills

- **\$skill-creator** - Bootstrap new skills from description
- **\$skill-installer** - Install skills from catalog
- **\$create-plan** (experimental) - Research and plan features

Install additional skills:

```
$skill-installer linear      # Linear integration  
$skill-installer notion     # Notion integration
```

Creating a Skill

```
# Use the built-in skill creator
$skill-creator Create a skill that generates
conventional commit messages based on staged changes
```

Codex will:

1. Create the skill folder structure
2. Generate SKILL.md with appropriate metadata
3. Add workflow instructions
4. Suggest reference files if needed

Skills vs Prompts

Aspect	Custom Prompts	Agent Skills
Location	<code>~/.codex/prompts/</code>	<code>~/.codex/skills/</code>
Structure	Single <code>.md</code> file	Folder with resources
Invocation	Slash commands	<code>\$skill-name</code> or auto
Resources	Text only	Scripts, templates, refs
Sharing	Copy files	Git-friendly folders

Recommendation: Use Skills for complex, multi-step workflows

Skills: Claude Code vs Codex

Aspect	Claude Code	Codex CLI
Spec	agentskills.io	agentskills.io
Format	SKILL.md + YAML	SKILL.md + YAML
User Location	<code>~/.claude/skills/</code>	<code>~/.codex/skills/</code>
Repo Location	<code>.claude/skills/</code>	<code>.codex/skills/</code>
Invocation	Auto/explicit	<code>\$skill-name</code> or auto
Creator	<code>skill-creator</code>	<code>\$skill-creator</code>

Same spec, different paths!

Configuration Profiles

Define Multiple Profiles

```
# ~/.codex/config.toml

[profiles.production]
model = "gpt-5.2-codex"
approval_policy = "on-request"
sandbox_mode = "workspace-write"

[profiles.development]
model = "gpt-5-codex-mini"
approval_policy = "never"
sandbox_mode = "danger-full-access"
```

Using Profiles

```
codex --profile production  
codex --profile development  
codex --profile testing
```


Resume Previous Sessions

```
# Open picker to choose a session
codex resume

# Resume the most recent session automatically
codex resume --last

# Resume a specific session by id
codex resume SESSION_ID
```

Session Lifecycle



Commands: `codex` , `codex resume` , `codex apply`

Session Commands

```
# Interactive session picker
codex resume

# Jump straight to the most recent session
codex resume --last

# Apply the last diff from the active session
codex apply
```

Non-Interactive Sessions


```
# Run in CI/CD pipeline
codex exec "update dependencies and fix breaking changes"

# Note: For resuming, use the regular command
codex resume
```

Advanced Features

Power User Tools

Unlock full potential

A blurred background image showing a computer screen with code snippets and a keyboard in the foreground.

```
link https://devel...  
package _s  
(! function_exists('incode_starter_setup')) {  
 * Sets up theme defaults and registers support for var-  
 * Note that this function is hooked into the init hook.  
 * runs before the init hook. The init hook is too late  
 * as indicating support for post thumbnails.  
 */  
function incode_starter_setup() {  
 /* Make theme available for translation.  
 * Translations can be filed in the /languages/ directory.  
 * If you're building a theme based on Twenty Sixty,  
 * use Twenty Sixty's build functions.  
 incode_starter_setup();  
}
```

Codex Cloud

Run long tasks without tying up your terminal:

```
# Launch a cloud task
codex cloud exec --env my-env "Refactor authentication module"

# Check task status
codex cloud status

# List running tasks
codex cloud list
```

Codex Cloud Benefits

- **Background execution** - Free up your terminal
- **Parallel tasks** - Run multiple jobs simultaneously
- **Persistent environments** - Pre-configured workspaces
- **Team collaboration** - Share environments and results
- **Long-running jobs** - Multi-hour refactoring sessions

IDE Extensions

Codex integrates with popular IDEs:

- **VS Code** - Full extension in marketplace
- **Cursor** - Native Codex support
- **Windsurf** - Integrated workflows

All extensions support:

- Skills and MCP servers
- Project-local configuration
- Same approval policies as CLI

Model Context Protocol

MCP Integration

Extend Codex with external tools



Model Context Protocol (MCP)

Configure MCP Servers

GitHub MCP Server

```
[mcp_servers.github]
command = "npx"
args = ["@modelcontextprotocol/server-github"]
env = { GITHUB_TOKEN = "${GITHUB_TOKEN}" }
```

Database MCP Server

```
[mcp_servers.postgres]
command = "npx"
args = ["@modelcontextprotocol/server-postgres"]
env = { CONNECTION_STRING = "${DATABASE_URL}" }
```

MCP Startup Guardrails (v0.31+)

Prevent flaky tools from freezing Codex:

```
[mcp_servers.github]
command = "npx"
args = ["@modelcontextprotocol/server-github"]
startup_timeout_ms = 15000 # Abort after 15 seconds
```

MCP Timeout Benefits

- Clean abort when helpers fail to boot
- Prevents entire run from freezing
- Better error messages
- Faster feedback on configuration issues

MCP Usage


- Automatically available when configured
- Access external tools and data
- Extend Codex capabilities

Running Codex as MCP Server

```
# Modern approach (v0.37+)
codex mcp --config ~/.codex/config.toml

# Legacy approach
codex serve --port 8080
```

MCP Architecture



MCP Server Benefits

The `codex mcp` command exposes Codex as a tool:

- Other agents can call Codex workflows
- IDEs can integrate without plugins
- Claude Code can use as sub-agent
- Mix model strengths (GPT + Claude)

Why Codex as Sub-Agent?

- Leverage GPT-5-Codex for complex tasks
- Use Codex's specialized prompts
- Access different model providers
- Unified approval/sandbox policies

Integration Example

```
# Add Codex as MCP server in Claude Code
claude mcp add codex -- codex mcp

# List MCP servers
claude mcp list

# Remove if needed
claude mcp remove codex
```

Integration Options

- Connect from other MCP clients
- Use in multi-agent workflows
- Integrate with IDEs

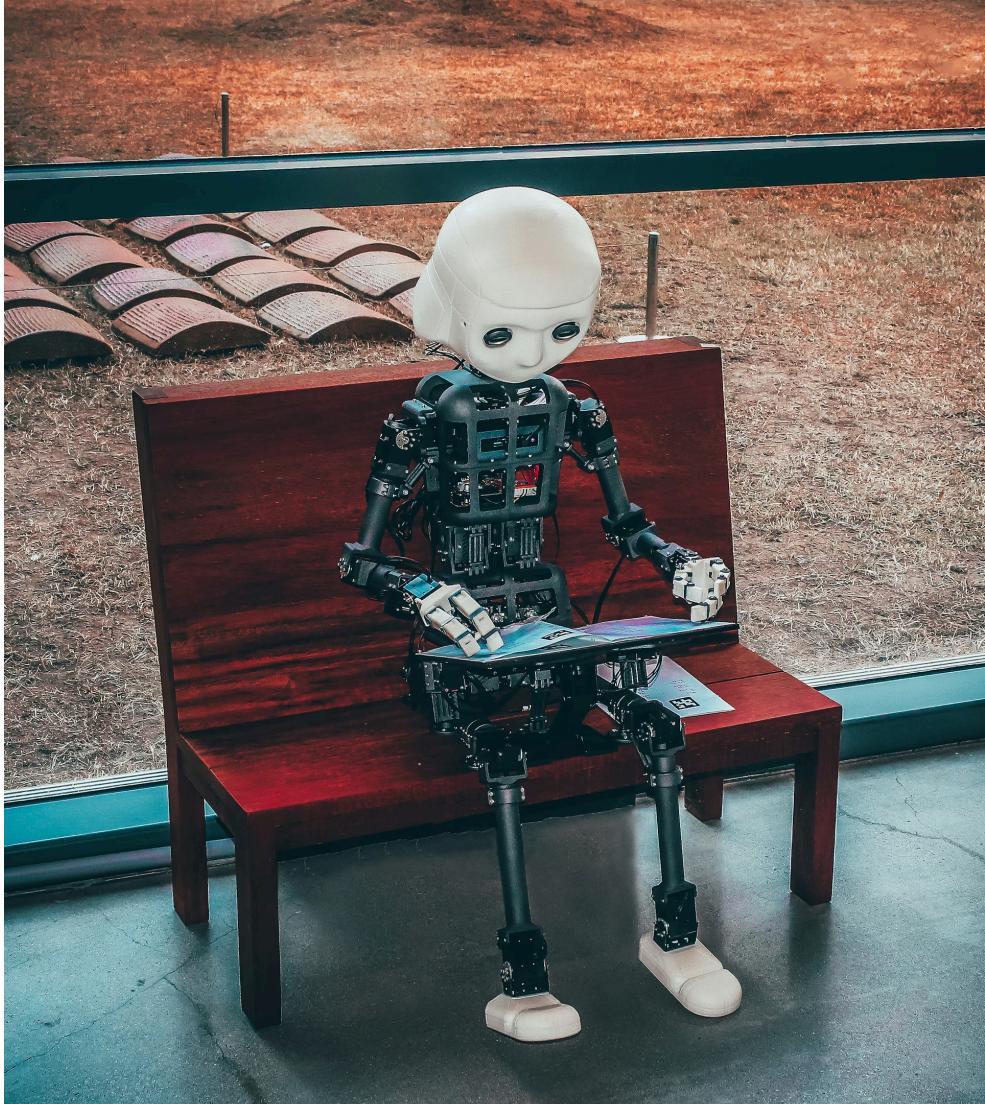
Example Client Connection

```
const client = new MCPClient({  
  url: 'http://localhost:8080',  
  capabilities: ['code-generation', 'review']  
});
```

Multi-Model Support

Provider Flexibility

OpenAI, Anthropic, Ollama, Azure



Multi-Model Provider Support

Configure Alternative Providers

Anthropic Claude

```
[providers.anthropic]
type = "anthropic"
api_key = "${ANTHROPIC_API_KEY}"
model = "claude-3-opus-20240229"
```

Local Ollama

```
[providers.ollama]
type = "ollama"
base_url = "http://localhost:11434"
model = "codellama"
```

Azure OpenAI

```
[providers.azure]
type = "azure-openai"
api_key = "${AZURE_API_KEY}"
endpoint = "https://myinstance.openai.azure.com"
deployment = "gpt-3.5-turbo"
```

Switch Providers

```
codex --provider anthropic  
codex --provider ollama  
codex --provider azure
```

Enable Detailed Logging

```
# Set log level
export RUST_LOG=debug
codex

# Trace level (maximum detail)
export RUST_LOG=trace
codex
```

Log Locations

- Interactive: `~/.codex/log/codex-tui.log`
- Non-interactive: stderr output
- Custom: Redirect with shell operators

Debug Configuration

```
# ~/.codex/config.toml
[logging]
level = "debug"
file = "/path/to/custom.log"
```

CI/CD & Automation

- GitHub Actions example lives in repo (`.github/workflows/codex-review.yml`)
- Typical steps: checkout → install Codex → authenticate → `codex exec` → upload artifacts
- Cron ideas: weekly security sweep, dependency refresh, monthly cleanup
- Guardrails: run on branches, review PRs before merge, notify on failures

CI/CD Pipeline Examples

```
# Fail-fast pipeline
git pull && \
codex exec "migrate database schema" && \
npm test
```

```
# Weekly cron example
0 2 * * 1 cd /path/to/repo && \
codex exec "weekly security audit"
```

```
# Chain commands to stop on failure
npm install && \
codex exec "fix any TypeScript errors" && \
npm run build
```

Advanced TOML Configuration

```
# ~/.codex/config.toml
model = "gpt-5.2-codex"
model_provider = "openai"
approval_policy = "on-request"
sandbox_mode = "workspace-write"
web_search_request = true
```

Environment & Notifications

```
[notification]
program = "notify-send"
args = ["Codex", "Task completed"]

[shell_environment]
NODE_ENV = "development"
PYTHONPATH = "/usr/local/lib/python3.9"
```

Shell Environment Policies

- **inherit** - Use parent shell environment
- **explicit** - Only specified variables
- **minimal** - Basic environment only

Environment Configuration

```
# Inherit all variables
shell_environment_policy = "inherit"

# Explicit variables only
shell_environment_policy = "explicit"
```

Explicit Environment

```
[shell_environment]
PATH = "/usr/local/bin:/usr/bin:/bin"
HOME = "/home/user"
LANG = "en_US.UTF-8"
```

Security Considerations

- Use `explicit` for production
- Use `inherit` for development
- Never expose secrets in config

Notification Options

- Notifications live in `[notification]` (see repo for full examples)
- macOS: `program="osascript"`, Linux: `program="notify-send"`, Webhook: `program="curl"`
- Triggers: task completion, approval prompts, error conditions (configurable)

Enterprise Features

- **MDM Configuration** - Managed settings on macOS via MDM profiles
- **Admin-scoped Skills** - Organization-wide skill deployment
- **Requirements.toml** - Enforce policies across teams
- **Zero Data Retention** - ZDR compliance with ChatGPT auth
- **Audit Logging** - Track all agent actions

Requirements.toml

Enforce organizational policies:

```
# /etc/codex/requirements.toml (UNIX)
# or via MDM (macOS)

[sandbox]
allowed_modes = ["read-only", "workspace-write"]
# Prevents danger-full-access

[skills]
admin_only = ["deploy-prod", "db-migrate"]

[approval]
required_for = ["file-delete", "git-push"]
```

Prompt Engineering Tips

When to Use Each Mode

Mode

Use Case

Interactive (`codex`)

Exploration, iteration, learning

Single prompt (`codex "..."`)

Quick questions, small tasks

Exec (`codex exec`)

Automation, CI/CD, scripts

Effective Prompts

Be specific about scope:

- "Fix the bugs"
- "Fix the null pointer exception in UserService.java line 42"

Provide context:

- "Add tests"
- "Add unit tests for the validateEmail function using Jest"

State expected outcomes:

- "Make it faster"
- "Optimize the database query to reduce response time below 100ms"

AGENTS.md Best Practices

Keep it focused and current:

```
# Project Context
E-commerce platform, Node.js + PostgreSQL

## Current Sprint
Payment integration with Stripe

## Conventions
- Use async/await, not callbacks
- All API responses follow { data, error } format
- Tests required for all new endpoints
```

Update AGENTS.md as your project evolves!



Practical Exercises

Hands-On Labs

Learn by doing

Exercise Structure

Three Main Categories

Available Labs

- **Lab 1:** Spring Boot REST API (generate from scratch)
- **Lab 2:** Python Refactoring (improve legacy code)
- **Lab 3:** React TypeScript Forms (frontend development)
- **Lab 4:** Microservices Architecture (multi-language)
- **Lab 5:** Skills Creation (extend Codex)

Each Exercise Includes

- Starter code (where applicable)
- Step-by-step Codex prompts
- Success criteria checklist
- Advanced challenges

Note: You build the solution using Codex—no reference implementations provided!

Lab 1: Spring Boot API

- Objective: Build a Spring Boot 3 task-management REST API end-to-end
- Timebox: 60–90 minutes
- Workspace: `exercises/java-spring-boot`
- Instructions: open `exercises/java-spring-boot/README.md`

Lab 2: Python Refactoring

- Objective: Modernize legacy Python code with clean architecture and tests
- Timebox: 45–60 minutes
- Workspace: `exercises/python-refactoring`
- Instructions: open `exercises/python-refactoring/README.md`

Lab 3: React TypeScript Forms

- Objective: Ship a production-ready registration flow with React, TypeScript, and Zod
- Timebox: 45–60 minutes
- Workspace: `exercises/react-forms`
- Instructions: open `exercises/react-forms/README.md`

Lab 4: Microservices

- Objective: Build an event-driven multi-language microservices system
- Timebox: 90–120 minutes
- Workspace: `exercises/microservices`
- Instructions: open `exercises/microservices/README.md`

Lab 5: Skills Creation

- Objective: Create a custom skill using `$skill-creator`
- Timebox: 30 minutes
- Task: Build a skill that generates conventional commit messages
- Workspace: `exercises/skills-creation/`
- Instructions: open `exercises/skills-creation/README.md`

Optional: Advanced Challenges

- **Database Migration:** Use Codex + MCP tools to modernize a legacy schema
- **AI Code Review:** Automate PR reviews using `.github/workflows/codex-review.yml`
- **Full-Stack Capstone:** Combine Labs 1-4 into a production-style application

These build on the core labs and can be explored as time permits.

Best Practices

Professional
Workflows

Enterprise-ready patterns



Review Changes Before Approving

- Codex displays unified diffs automatically
- Use `/diff` to see all pending changes
- Review line-by-line for unintended edits
- Check file statistics (insertions/deletions)
- Catch mistakes before they land in codebase

Pro tip: Always review diffs for:

- Accidental deletions
- Unrelated file changes
- Security implications

Security Best Practices

Sandbox Configuration by Environment

Development Profile

```
[profiles.dev]
sandbox_mode = "danger-full-access"
approval_policy = "never"
```

Staging Profile

```
[profiles.staging]
sandbox_mode = "workspace-write"
approval_policy = "on-request"
```

Production Profile

```
[profiles.prod]
sandbox_mode = "read-only"
approval_policy = "on-request"
```

Security Guidelines

- Never store API keys in config files
- Use environment variables for secrets
- Enable approval for production
- Regular audit of generated code
- Restrict network access in sandbox

Team Collaboration

Shared Configuration

Project AGENTS.md

```
# Team: Platform Engineering
## Conventions
- PR reviews required for all changes
- Follow company style guide
- Security scanning mandatory
- 80% test coverage minimum
```

Current Sprint Context

```
## Current Sprint
- Migrating to Kubernetes
- Implementing OAuth 2.0
```

Shared Prompts Repository

```
# Clone team prompts
git clone team-repo/codex-prompts ~/.codex/prompts

# Keep synchronized
cd ~/.codex/prompts && git pull
```

Model Selection Strategy

```
[profiles.quick]
model = "gpt-5-codex-mini" # Fast responses
```

```
[profiles.complex]
model = "gpt-5.2-codex" # Complex reasoning
```

Local Models

```
[profiles.local]
provider = "ollama"
model = "codellama" # No API costs
```

Optimization Tips

- Use smaller models for simple tasks
- Cache responses with session resumption
- Batch similar operations
- Use local models for sensitive data

Troubleshooting

Common Issues

Solutions and workarounds



Troubleshooting Guide

Common Issues & Solutions

Authentication Failures

```
# Clear cached credentials
rm -rf ~/.codex/auth
```

```
# Re-authenticate
codex login --headless
```

Sandbox Errors

```
# Confirm sandbox configuration
grep sandbox_mode ~/.codex/config.toml

# Bypass for Docker environments
codex --dangerously-bypass-approvals-and-sandbox
```

MCP Connection Issues

```
# Test MCP server
npx @modelcontextprotocol/server-test

# Enable debug logging
RUST_LOG=trace codex
```

Context Limit Errors

Symptoms: "Context window exceeded", slow responses

Solutions:

- Start a new session: `/clear` or new terminal
- Use `codex resume` to continue with trimmed context
- Break large tasks into smaller prompts
- Remove verbose files from AGENTS.md

Model Availability Issues

```
# Check available models
codex --model list

# Fall back to different model
codex --model gpt-5-codex-mini

# Verify API connectivity
curl -I https://api.openai.com/v1/models
```

Common Error Messages

Error	Cause	Fix
"Token expired"	Auth timeout	<code>codex login</code> again
"Rate limited"	Too many requests	Wait or use <code>--model mini</code>
"No such table"	Wrong osquery platform	Check platform docs
"Sandbox denied"	Permission blocked	Adjust sandbox mode

VS Code Integration

```
{  
  "tasks": [ {  
    "label": "Codex Review",  
    "type": "shell",  
    "command": "codex exec 'Review ${file} for issues'"  
  }]  
}
```

Git Hooks

```
#!/bin/bash
# .git/hooks/pre-commit
codex -n --profile review \
    "Check staged files for security issues"
```

Make Integration

```
review:
```

```
    codex exec "Review all changes since last commit"
```

```
generate-tests:
```

```
    codex exec "Generate missing unit tests"
```

Structured Logging

```
[logging]
level = "info"
format = "json"
file = "/var/log/codex/codex.log"
```

Metrics Configuration

```
[metrics]
enable = true
endpoint = "http://metrics.internal:9090"
```

Log Analysis

```
# Parse JSON logs
cat ~/.codex/log/codex-tui.log | jq '.level == "error"'

# Monitor in real-time
tail -f ~/.codex/log/codex-tui.log | grep ERROR
```

Build Your Own MCP Server

```
// custom-mcp-server.js
import { MCPServer } from '@modelcontextprotocol/sdk';

const server = new MCPServer({
  name: 'custom-tools',
  version: '1.0.0'
});
```

MCP Tool Definition

```
tools: [{

  name: 'database-query',
  description: 'Execute database queries',
  handler: async (params) => {
    return { result: 'Query executed' };
  }
}]
```

Register Custom Server

```
[mcp_servers.custom]
command = "node"
args = ["./custom-mcp-server.js"]
```

From GitHub Copilot

- Export commonly used snippets
- Convert to Codex prompts
- Leverage session persistence

From Claude Code

```
# Import Claude memory
cp ./claude-code/.claude/CLAUDE.md ./AGENTS.md

# Convert slash commands
for f in ./claude-code/.claude/commands/*.md; do
    cp "$f" ~/.codex/prompts/
done
```

From Cursor/Codeium

- Migrate project context
- Recreate custom instructions
- Set up equivalent workflows

Recent Features (v0.40-0.87)

Version 0.40-0.87 Highlights

- **Agent Skills** (v0.76) - Reusable instruction bundles
- **GPT-5.2-Codex** (v0.81) - New default API model
- **Project-local config** (v0.78) - `.codex/config.toml` per-repo
- **Multi-agent control** (v0.79) - Spawn/message conversations
- **Thread rollback** (v0.79) - Undo last N turns
- **Elevated sandbox** (v0.80) - `/elevate-sandbox` command
- **Ctrl+G editor** (v0.78) - Open prompt in external editor
- **Enterprise MDM** (v0.78) - Managed configuration on macOS

Earlier Features (v0.30-0.39)

- MCP startup timeouts (v0.31)
- Token usage in /status command (v0.35)
- GPT-5-Codex high reasoning mode (v0.36)
- Network allowlists for testing (v0.36)
- Simplified MCP server: `codex mcp` (v0.37)
- Hierarchical AGENTS.md cascading (v0.39)

MCP Robustness (v0.31+)

Always add timeout to MCP servers:

```
[mcp_servers.your_server]
command = "your-command"
startup_timeout_ms = 15000 # Recommended
```

Advanced Model Control (v0.36+)

Use GPT-5-Codex for complex, long-running tasks:

```
codex -m gpt-5-codex -c model_reasoning_effort='high'
```

Or configure in TOML:

```
model_reasoning_effort = "high" # minimal/low/medium/high
```

High Reasoning Mode Benefits

- Multi-hour work sessions allowed
- Iterates tests until green
- Deep problem-solving capability
- Automatic retry on failures
- Best for complex refactoring

When to Use High Reasoning

- Large test suite fixes
- Complex architectural changes
- Multi-file refactoring
- Performance optimizations
- Breaking change migrations

Network Access Control

Control network access in sandbox mode:

```
# ~/.codex/config.toml
[sandbox_workspace_write]
network_access = true # Default: false
```

Note: Domain-specific allowlists may be available in v0.36+ (check release notes)

Network Control Benefits

- Offline by default for reproducibility
- Allow only specific staging APIs
- Prevent accidental external calls
- Maintain test isolation
- Control data exfiltration

Network Allowlist Use Cases

- Integration tests with staging APIs
- CI/CD pipelines with controlled access
- Development with specific endpoints
- Security-sensitive environments
- Reproducible test suites

What Shipped in 2025

- **Agent Skills** - Reusable instruction bundles (Dec 2025)
- **VS Code Extension** - IDE integration shipped
- **Codex Cloud** - Launch cloud tasks from CLI
- **Multi-agent coordination** - Thread spawning and messaging
- **GPT-5 Model Family** - 5.2-Codex, 5.1-Codex-Max, Mini

What's Coming (2026)

- Enhanced skills marketplace
- Deeper IDE integrations
- Advanced collaboration features
- Extended platform support

Community & Ecosystem

- Open source at github.com/openai/codex
- 56k+ GitHub stars, 400+ releases
- Skills catalog at github.com/openai/skills
- Active discussions and contributions welcome

Essential Commands

```
# Basic usage
codex                         # Interactive mode
codex exec "prompt"           # Execute task & exit
codex resume                   # Resume session
codex apply                     # Apply last diff
```

Configuration Commands

```
codex --profile production          # Use profile  
codex --sandbox read-only          # Set sandbox  
codex --ask-for-approval on-request # Set approval
```

Advanced Commands

```
codex mcp --config ~/.codex/config.toml # MCP server mode (v0.37+)
codex apply                         # Apply last diff
codex resume --last                  # Resume most recent session
```

Key Files

- Config: `~/.codex/config.toml`
- Prompts: `~/.codex/prompts/*.md`
- Memory: `./AGENTS.md`
- Logs: `~/.codex/log/`

Documentation & Code

Official Documentation

<https://github.com/openai/codex/docs>

GitHub Repository

<https://github.com/openai/codex>

Course & Community

Course Materials & Labs

<https://github.com/kousen/codex-training>

Community Support

<https://github.com/openai/codex/discussions>

MCP Registry

<https://modelcontextprotocol.io/registry>

Discord & Office Hours

Discord Server

Join the Codex community for support and discussion

Weekly Office Hours

Every Thursday at 2 PM PST

Contributing

- Bug reports: GitHub Issues
- Feature requests: GitHub Discussions
- Code contributions: Pull Requests

Ecosystem

- MCP server templates
- Prompt libraries
- Configuration examples
- Integration guides

Codex CLI vs Claude Code

For teams using both tools:

Aspect	Codex CLI	Claude Code
Model	GPT-5.2-Codex	Claude Opus 4.5
Runtime	Rust	TypeScript
Memory File	AGENTS.md	CLAUDE.md
Skills Location	~/.codex/skills/	~/.claude/skills/
Skill Invoke	\$skill-name	Auto or explicit
Cloud Tasks	Codex Cloud	-

When to Use Which?

Choose Codex CLI when:

- GPT-5 model strengths needed
- Codex Cloud for long tasks
- Enterprise MDM requirements

Choose Claude Code when:

- Claude's reasoning preferred
- Extended thinking needed
- Anthropic ecosystem integration

Use both with MCP to leverage each model's strengths!

References & Credits

Newsletter Sources

MLearning.ai Art on Substack by @mlearning

- [100 OpenAI Codex CLI Tricks and Tips](#)
- [30 Codex CLI Tips v0.30-0.39](#)

Many advanced tips including:

- Command chaining with `&&`
- Scheduled maintenance automation
- Network allowlists
- High reasoning mode

Additional Resources

Official Sources

- [Codex CLI Documentation](#)
- [Agent Skills Guide](#)
- [Codex GitHub Repository](#)
- [Skills Catalog](#)

Specifications & Community

- [agentskills.io - Skills specification](#)
- [Model Context Protocol](#)
- [Codex Discussions](#)

Related Training

- [Claude Code Training](#)

Thank You!

Questions?



Kenneth Kousen*Author, Speaker, Java & AI Expert*

kousenit.com | [@kenkousen](https://twitter.com/kenkousen)